

Course Project Title

Social Network Analysis for Computer Scientists — Course Project Paper

Sevak Mardirosian
LIACS, Leiden University
s.mardirosian@umail.leidenuniv.nl

Michail Tsiaousis
LIACS, Leiden University
m.tsiaousis@umail.leidenuniv.nl

ABSTRACT

Abstract.

Categories and Subject Descriptors

H.4 [Information Systems Applications]: Miscellaneous;

D.2.8 [Software Engineering]: Metrics—*complexity measures, performance measures*

General Terms

Theory

Keywords

ACM proceedings, L^AT_EX, text tagging

1. INTRODUCTION

One of the main ideas behind link prediction is that it offers several ways in which one can study networks of any kind in which data scientists, software engineers and researchers could greatly benefit from. Social networks are a popular way to model the interactions among people in a group or community. Their connections can be visualized as a graph, where a vertex corresponds to a person in a group and an edge represents some form of association between the corresponding persons.

Organizations, such as Amazon, can extract information based on raw data in order to predict what the customer might actually buy or find interesting. Another example would be the professional social-media platform, Linked-in. Based on incoming data they would be able to for instance, predict your next connection; link back jobs which are relevant to you; link contents/articles based on your connections, interests, and things you read on the web. Similarly, security corporate companies could more precisely focus their efforts based on probable relationships in malicious networks that have heretofore gone unobserved, and researchers can easily adapt link prediction methods to iden-

This paper is the result of a student course project, and is based on methods and techniques suggested in [10, 14]. Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice on the first page. SNACS '17 Social Network Analysis for Computer Scientists, Master CS, Leiden University (liacs.leidenuniv.nl/~takesfw/SNACS).

tify links that are surprising given their surrounding network, or links that may not exist at all [14].

That being said, links or associations are usually based on someone's behavior and interests. In social networks however, objects are dynamic given that the number of edges and nodes are changing continuously. Getting a hold of those dynamics [number of nodes and edges] is one of the fundamental problems in social network analysis. The main problem which we are going to cover in this paper is, given a network at time t , accurately predict the links that are going to form at a future time t' . This is known as the link prediction problem.

In this paper, we approach the problem of link prediction using supervised learning [8, 10, 14, 22]. Unsupervised techniques have also been used to solve the problem of link prediction. The main difference between the two is that in the unsupervised case, we calculate metrics that individually provide predictions of whether a link will form or not. Supervised learning employs a more sophisticated framework where each of these metrics could be considered as an input to a machine learning model.

This paper is based on methods and techniques described in [10] and [14]. Both papers consider collaboration networks on which several machine learning models are applied in order to predict future links (collaborations) between authors.

In the framework of supervised learning, the prediction task is framed as a binary classification problem where a classifier is built in order to distinguish between the instances of the positive and negative class. We say that a pair of edges belongs to the positive class if they are connected, and to the negative class otherwise. One of the main problems in link prediction is that of class imbalance, that is the number of instances of the negative class are overwhelmingly greater than those of the positive. This means that the number of links that do actually form (positive class) is dramatically less than the number of links that do not form (negative class) [14]. The reason that class imbalance consists a problem is due to the fact that it impedes the ability of classifiers to learn from the minority class [2], resulting to a classifier that predicts every instance as belonging to the majority class.

Another method for dealing with imbalanced data sets is that of applying a specific type of decision tree called Hellinger

Distance Decision Tree (HDDT) [1, 2], which are robust and insensitive to skewed distributions. This method is suggested in [14] as a solution to class imbalance in link prediction, but the authors have not provided results of this method. To the best of our knowledge, HDDT have not been applied to solve the link prediction problem. In this paper, we compare C4.5 decision trees [19] and HDDT as applied and proposed in [10] and [14], respectively. Our interest is to examine whether or not HDDT provide an alternative solution to the link prediction problem, in which class imbalance is an important issue. Refer to Section 4 for an analysis of HDDT and C4.5.

The paper is structured as follows. In Section 2, we formalize the link prediction problem. Specifically, we discuss the transformation of a network to a tabular data set, and the class imbalance problem. Related work is provided in Section 3. In Section 4, we describe the approaches that have been taken in [10, 14], and we also provide an analysis of HDDT and C4.5. Section 5 describes the data sets and software in which the experiments and results of Section 6 are based on. Finally, we conclude in Section 7.

2. PROBLEM STATEMENT

2.1 Preliminaries

Consider at a particular time t , an undirected network $G = \langle V, E \rangle$, where V is the set of nodes and E is the set of edges. The goal is to predict the edges that will form in the network at time $t' > t$. As an example, consider a collaboration network of authors. Two authors v and w are linked when they have published at least one research paper together. The goal is to predict future connections between authors that have not yet collaborated.

In this paper, we approach the problem of link prediction using supervised learning. In this framework, in order to build models with predictive ability, one needs to construct two sets, mainly a *train* and a *test* set. Then, the model learns features and patterns from the train set, and it is evaluated on the test set. These two sets form together the *data set*. This set is of the form (\vec{x}, y) , where \vec{x} indicates the inputs to the model, and y indicates the label, that is whether or not a link is present between two nodes. The model learns from the instances (\vec{x}, y) of the train set. In the evaluation phase, we feed the instances \vec{x} of the test set to the trained model, and the output is the predicted label, i.e. our prediction of whether or not a link will form between two nodes.

2.2 Constructing Data Sets

The process of constructing the train and test set is described in detail in [10, 14]. In order to transform the network to a data set, we first choose two adjacent periods, the *train* and *test period*. In the train period, denoted by $[0, t_k]$, we have pairs of nodes that are connected, that is, they have some sort of association or connection. The same holds for the test period $[t_{k+1}, t_l]$. We construct a new network $G_k = \langle V_k, E_k \rangle$ consisted of all pair of nodes that do not share a connection in the train period. We then examine whether or not they are connected in the test period. If they did, we assign to the pair a positive label, and otherwise a negative label. The data set contains $\binom{|V_k|}{2} - |E_k|$ instances, since we do not account for self-loops, and in undirected networks each pair is recorded once.

Selecting an appropriate set of features is one of the most important steps in supervised learning algorithms [10]. The inputs to the model are selected by calculating topology-based and node-based metrics from G_k that will serve as attributes to the data set. Such metrics are, but not limited to, *common neighbors*, *Jaccard coefficient*, *Adamic Adar*, *preferential attachment*, *Katz measure*, *Rooted PageRank*, and many others. All these features can be used individually to provide naive predictions. This is sometimes called the unsupervised framework of link prediction. In the supervised case, we use state-of-the-art predictive models such as Neural Networks, Support Vector Machines, Decision Trees, and Random Forests that treat these individual metrics as inputs. The output or prediction of the model, is whether or not a link will form between two specific nodes.

That being said, we end up with a data set in the standard format (\vec{x}, y) , where \vec{x} are metrics extracted from G_k and y indicates the label, that is whether or not a pair of nodes is connected or not.

2.3 Class Imbalance and Size

Although we have constructed the data set, there are two main problems. The first problem is that its size is very large which can make the training process infeasible. For example, consider the relatively small in size **MathOverflow** network [17] that can be found in the SNAP repository¹, consisted of 21688 nodes and 107581 edges. Defining the train period to be two-thirds of the networks' edges, and assuming that all node appear in V_k , the resulting data set would be of size $\binom{21688}{2} - 107581$, which yields more than 235 million instances.

As we mentioned in Section 1, one of the main problems in link prediction is the imbalance between the positive and negative class. That is, the amount of instances in the constructed data set with a negative label is overwhelmingly greater than those with a positive label.

A technique called under-sampling provides a solution to the class imbalance problem. In under-sampling, we acquire a representative sample from the majority class that is more or less equal to the size of the minority class. In this way, we create a balanced data set comprised from the sample of the majority class and all instances of the minority class. We then say that the data set is under-sampled to balance. Note that, when a data set is under-sampled to balance means that the ratio of positive and negative instances is 1:1. It is possible to under-sample the data set at different ratios. To see this, consider the case where the class imbalance is of size 1:20000 and the data set has 10 million instances. Training a model with this setup is a difficult task due to the huge size of the data set. In this case, it would be helpful to under-sample the data set to a ratio different than 1:1, for example to a ratio of 1:10. That is, we explicitly acquire a sample of the majority class that is ten times larger than that of minority's class. In this way, size of the data set is decreased and at the same time we retain imbalanced classes.

A second method that mitigates both problems of size and class imbalance is proposed in [14]. As we mentioned in

¹<http://snap.stanford.edu/index.html>

Section 2.2, in the process of constructing the data set we consider each pair of nodes that do not exist in the train period as one that could potentially connect in the future. It is suggested to treat each neighborhood as a separate problem. We define the n -neighborhood of a node v as the set of nodes that are 0 to n steps away from v . Instead of considering all pairs of nodes to enter the data set, we choose for each node, and for a specified n , only those that belong to the n -neighborhood of the node. In this way, both the size of the data set and the imbalance between the classes decrease greatly. For the case of **MathOverflow** network, choosing $n = 2$ results to a data set with more than 5 million instances of which only 3168 belong to the positive class, an imbalance of about 1:1695.

3. RELATED WORK

The link prediction problem was formalized in [12], where there is an extensive analysis of different baseline (naive) predictors, applied to co-authorship networks. Nevertheless, state-of-the-art supervised learning models are not studied in this paper. Furthermore, co-authorship networks have been analyzed in [10], [18], and [4], but this time several machine learning algorithms such as Decision Trees, SVM, Neural Networks, and K-Nearest Neighbors are used for prediction. In contrast to [10], where static, unweighted networks are used, [18] takes into account the evolution of the network, and the authors of [4] analyze a weighted network. Several recommendations and suggestions are presented in [14], regarding how to perform link prediction, and how to approach inherent issues such as dealing with class imbalance, variance reduction, and huge train sets.

In the domain of counter-terrorism, link predictions has been studied in [5], [8]. In the latter paper, several networks are constructed by removing nodes from the original graph, called visible networks. A set of features is calculated from these networks in order to predict links of the original one, using supervised learning models. A more probabilistic approach is taken in [6], where chance-constrain programs are used, and in [21], in which the framework of Relational Markov Network is utilized. Finally, [23], and [11] are extensive surveys regarding the different metrics, classification methods, and approaches that have been used in the link prediction domain.

4. SUGGESTED APPROACHES

In this section, we describe the approaches that have been taken in [10] and [14] and we discuss the differences between them.

4.1 Construction of Data Sets

The authors in [10] use two undirected, unweighted collaboration networks of co-authors, mainly the BIOBASE² and DBLP³ networks. An edge between two authors/nodes exists in case they have written at least on research paper together. Both networks contain timestamps of the years of collaborations. For each network, the authors choose a corresponding train and test period that span one or more years. For example, in the first network, the train period is consisted of 5 years from 1998 to 2002, whereas the test

period comprises the year 2003. The data set constructed from this network contains each combination of co-authors that do not exist in the train period, that is, each pair that could potentially have a connection in the future. For each of these combinations, it is examined whether or not there is a connection in the test period. If a connection exists, a positive label is assigned to the pair, otherwise a negative label.

In [14], the authors use two networks, one weighted, undirected co-authorship network of 19464 edges of condensed matter physics collaborations from 1995 to 2000 called **cond-mat**, and one directed weighted network of 712 million cellular phone calls, called **phone**. The process of constructing the data sets is restricted in the n -neighborhood of each node. Specifically, instead of considering all possible combinations of pairs of nodes that do not exist in the train period, the authors construct the data set only from those combinations that are not part of the train period, but also contained in the neighborhood of n size. Specifically, for each node v , the n -size neighborhood is calculated. We then compute the metrics between v and the nodes of this neighborhood, by excluding pairs that already belong to the train period. The idea behind this is that, it is more possible for links to form in the future between nodes that are closely together, than with nodes that are far apart. Thus, reducing the problem to each neighborhood offers two advantages. Firstly, the imbalance between classes is decreased, and secondly given that specific metrics such as *Adamic Adar* and *Rooted PageRank* are expensive to compute, we avoid calculating those metrics for pairs of nodes that are not highly likely to connect in the future. The value of n is chosen by the user. Selecting higher values guarantee that we do not lose connections that actually do form, but on the other hand the instances of the negative class increases dramatically, along with the computation time of the metrics.

4.2 Calculation of Features

The authors in [10] calculate several topological and domain specific features. Topological features include the *shortest distance* and *clustering index*, whereas the domain specific features include the sum of papers the two authors share, the sum of the count of keywords in the papers of each author, the number of common keywords two authors have used etc.

On the other hand, the authors of [14] concentrate on features that can be calculated for every network. The High Performance Link Prediction (HPLP) framework is presented which comprises several topological metrics such as *In/Out Degree*, *Common Neighbors*, *Maximum Flow*, *Adamic Adar*, *Jaccard Coefficient*, and others. Note that this is different from the work of Hasan et al. [10] where most of the selected features are domain specific. Hence, HPLP serves as a general framework for link prediction. Additionally, a new unsupervised prediction method is presented called **PropFlow**. It is a predictor based on random walk that starts from node v_i and ends at a node v_j in l or fewer steps, using link weights as transition probabilities [14]. A score s_{ij} is produced for each pair of nodes v_i and v_j that serve as an estimation of the likelihood that the two nodes will be connected in the future. **PropFlow** is also included in HPLP.

4.3 Overcoming Imbalance

²<http://www.elsevier.com>

³<http://dblp.uni-trier.de/xml/>

In [10] the problems of huge training sets and that of class imbalance are solved by under sampling the data set to balance. After creating a balanced data set, the authors train several machine learning models such as Decision Trees (with and without Bagging), Support Vector Machines, Naive Bayes, Neural Networks and others. They found that Bagging and Support Vector Machines provided the best performances when compared to different metrics such as accuracy, precision, recall, F-value and squared error. The models are evaluated using 5-fold cross validation.

Under-sampling to balance is also used in [14]. The authors use the features proposed in the HPLP framework and apply Random Forests with bagging, for different values of the neighborhood size n . They found that with increasing values of n there is a deterioration of performance in both networks. The value of $n = 2$ provided the optimal results.

4.4 HDDT and C4.5

In this sub-section we give a brief introduction to decision trees and we discuss the different ways in which C4.5 and HDDT are constructed.

Classification trees are tree-like structures that consist of nodes that can be either a root node, internal nodes, or terminal nodes, also called leaves. The nodes are connected with each other from top to bottom with edges, also called branches. The root node is the first node of the tree, while the terminal nodes(leaves) are the final nodes. For each node, there are two outgoing branches, the *left* and *right* branch that lead to the next two nodes. In order to classify an instance to one of the two classes, in each internal node, we test whether or not a specified condition is satisfied. If it is, we follow the left branch, and otherwise the right one. By evaluating over those conditions we arrive at the leaves that indicate the class prediction for the particular instance.

4.4.1 Recursive Partitioning

Consider n observations of p features X_1, X_2, \dots, X_p , and a class variable Y that indicates the class for each instance. In our case, n is the number of instances in the data set, $X_i, i = 1, 2, \dots, p$ indicate the calculated features from the network G_k of the train period, and variable Y corresponds to the class, that is whether or not an instance has formed in the test period. Variables $X_i, i = 1, 2, \dots, p$ form a p dimensional space in which the observations lie. Decision trees recursively partition this space into non-overlapping rectangles. The partitioning is accomplished in the following way. One of the X_i variables is selected, say X_r , and splits the p -dimensional rectangle into two parts. One part contains all observations for which $X_r \leq s_r$, and the second part contains all observations for which $X_r > s_r$, where s_r is a value of feature X_r . The remaining parts are partitioned in the same fashion by either the same variable or an other $X_i, i = 1, 2, \dots, p, p \neq r$. The goal is to split this hyper dimensional space in the best possible way, that is partition the space such as each rectangle contains observations only from a specific class.

Figure 1 taken from [15], illustrates the partitioning of 2-dimensional space with features X_1, X_2 on the left and the corresponding decision tree on the right. The tree has two internal nodes and four leaves. The condition of the root




Figure 1: text

node tests for each instance of the data set, whether or not $X_2 \leq 0.7$. If the condition holds we follow the left branch where we find the second condition, this time with respect to X_1 . If $X_1 \leq -1.4$ holds we predict that the instance belongs to the red class, otherwise to the green class.

4.4.2 Splitting Criteria

5. DATA SETS AND SOFTWARE TOOLS

5.1 Data Sources

Almost all link prediction works need to verify their methods on the collected datasets. The datasets are important for fairly reproducing and comparing different link prediction methods. Constructing and collecting the datasets is a time-consuming and labor-intensive work. That said, not all the datasets are publicly available to use and some of them are incomplete. Since the process of transforming the network to a data set requires the creation of two non-overlapping periods that simulate the formation of links between nodes from one period to the other, it is essential to consider networks that contain timestamps. We are interested in both directed and undirected networks. Nevertheless, we convert the former type to the latter. In all of our experiments, we consider unweighted networks.

A summary of the size of nodes and edges, along with the length of the timespan for the 5 social networks we consider can be found in Table 1. We use three networks that come from the KONECT⁴ database. The first one is the UC Irvine⁵ [16], a directed network containing sent messages between the users of an on-line community of students from the University of California, Irvine. Next, Digg⁶ [3] is a reply directed network of the social news website Digg where each node in the network is a user of the website, and each directed edge denotes that a user replied to another user. Furthermore, we use the Slashdot⁷ [9], a directed reply network of the technology website Slashdot, where nodes and edges represent users and replies, respectively. Moreover, we consider the RealityCall⁸ [7, 20] an undirected

⁴<http://konect.uni-koblenz.de/networks/>

⁵<http://konect.uni-koblenz.de/networks/opsahl-ucsocial>

⁶http://konect.uni-koblenz.de/networks/munmun_digg_reply

⁷<http://konect.uni-koblenz.de/networks/slashdot-threads>

⁸<http://networkrepository.com/ia-reality-call.php>

	UC	Digg	Slashdot	Reality	MathOver
Nodes	1899	30398	51083	6809	21688
Edges	59835	87627	140778	7680	107581
Timespan	6 months	16 days	2 years	4 months	6.5 years

Table 1: Number of nodes, edges, and length of timespan for the 5 social networks.

network that consists of human mobile phone call events between a small set of core users at the Massachusetts Institute of Technology (MIT) whom actually were assigned mobile phones for which all calls were collected. The network also contains edges between users that do not belong in the small set of users, who called other individuals that were not actively monitored. Therefore, some sort of noise might exist in this network, since we do not have more information about the individuals outside of the small set of MIT users. Finally, we consider the directed **MathOverflow**⁹ [17] network, where a node represents a user, and an edge indicates that user u answered user's v question on the website. All five networks are well suited for the link prediction problem. This is due to the fact that the nature of these networks is the interaction of individuals and we expect that the current interactions form a base for the connections that will form in the near future. Hence, we learn the patterns from interactions taking place at time $[0, t]$, in order to predict future connections for $t' > t$.

5.2 Software Tools

Although there are many link prediction metrics and methods proposed, only very few works open their source codes. Re-implementing methods and formulas to calculate predictors is a time-consuming process. Only few public tools try to integrate these metrics and methods such as **linkpred**¹⁰ and **LPmade**¹¹ [13] which both have a handful of link prediction metrics. **LPmade** is a cross-platform software solution that provides multi-core link prediction and related tasks and analysis [13]. It is written in **C++** and therefore it is suited for handling very large networks. Unfortunately, compilation issues and furthermore the lack of documentation and support prevented us from using the software. Hence, we used a **Python** library called **linkpred**. The main disadvantage of **linkpred** is that it is entirely written **Python** and it can prove to be very slow for handling large networks.

6. REFERENCES

- [1] D. A. Cieslak and N. V. Chawla. *Learning Decision Trees for Unbalanced Data*, pages 241–256. Springer Berlin Heidelberg, Berlin, Heidelberg, 2008.
- [2] D. A. Cieslak, T. R. Hoens, N. V. Chawla, and W. P. Kegelmeyer. Hellinger distance decision trees are robust and skew-insensitive. *Data Mining and Knowledge Discovery*, 24(1):136–158, Jan 2012.
- [3] M. De Choudhury, H. Sundaram, A. John, and D. D. Seligmann. Social synchrony: Predicting mimicry of user actions in online social media. In *Computational Science and Engineering, 2009. CSE'09. International Conference on*, volume 4, pages 151–158. IEEE, 2009.
- [4] H. R. De Sá and R. B. Prudêncio. Supervised link prediction in weighted networks. In *Neural Networks (IJCNN), The 2011 International Joint Conference on*, pages 2281–2288. IEEE, 2011.
- [5] M. Dombroski, P. Fischbeck, K. Carley, et al. Estimating the shape of covert networks. In *Proceedings of the 8th International Command and Control Research and Technology Symposium*, 2003.
- [6] J. Doppa, J. Yu, P. Tadepalli, and L. Getoor. Chance-constrained programs for link prediction. In *NIPS Workshop on Analyzing Networks and Learning with Graphs*, 2009.
- [7] N. Eagle and A. Pentland. Reality mining: sensing complex social systems. *Personal and Ubiquitous Computing*, 10(4):255–268, 2006.
- [8] M. Fire, R. Puzis, and Y. Elovici. *Link Prediction in Highly Fractional Data Sets*, pages 283–300. Springer New York, New York, NY, 2013.
- [9] V. Gómez, A. Kaltenbrunner, and V. López. Statistical analysis of the social network and discussion threads in slashdot. In *Proceedings of the 17th international conference on World Wide Web*, pages 645–654. ACM, 2008.
- [10] M. A. Hasan, V. Chaoji, S. Salem, and M. Zaki. Link prediction using supervised learning. In *In Proc. of SDM 06 workshop on Link Analysis, Counterterrorism and Security*, 2006.
- [11] M. A. Hasan and M. J. Zaki. *A Survey of Link Prediction in Social Networks*, pages 243–275. Springer US, Boston, MA, 2011.
- [12] D. Liben-Nowell and J. Kleinberg. The link prediction problem for social networks. In *Proceedings of the Twelfth International Conference on Information and Knowledge Management, CIKM '03*, pages 556–559, New York, NY, USA, 2003. ACM.
- [13] R. N. Lichtenwalter and N. V. Chawla. Lpmade: Link prediction made easy. *Journal of Machine Learning Research*, 12(Aug):2489–2492, 2011.
- [14] R. N. Lichtenwalter, J. T. Lussier, and N. V. Chawla. New perspectives and methods in link prediction. In *Proceedings of the 16th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, KDD '10*, pages 243–252, New York, NY, USA, 2010. ACM.
- [15] W.-Y. Loh. Classification and regression trees. *Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery*, 1(1):14–23, 2011.
- [16] T. Opsahl and P. Panzarasa. Clustering in weighted networks. *Social networks*, 31(2):155–163, 2009.
- [17] A. Paranjape, A. R. Benson, and J. Leskovec. Motifs in temporal networks. In *Proceedings of the Tenth ACM International Conference on Web Search and Data Mining, WSDM '17*, pages 601–610, New York, NY, USA, 2017. ACM.
- [18] M. Pavlov and R. Ichise. Finding experts by link prediction in co-authorship networks. In *Proceedings of the 2Nd International Conference on Finding Experts on the Web with Semantics - Volume 290, FEWS'07*, pages 42–55, Aachen, Germany, Germany, 2007. CEUR-WS.org.
- [19] J. R. Quinlan. *C4. 5: programs for machine learning*. Elsevier, 2014.

⁹<http://snap.stanford.edu/data/sx-mathoverflow.html>

¹⁰<https://github.com/rafguns/linkpred>

¹¹<https://github.com/rlichtenwalter/LPmade>

- [20] R. A. Rossi and N. K. Ahmed. The network data repository with interactive graph analytics and visualization. In *Proceedings of the Twenty-Ninth AAAI Conference on Artificial Intelligence*, 2015.
- [21] B. Taskar, M. fai Wong, P. Abbeel, and D. Koller. Link prediction in relational data. In S. Thrun, L. K. Saul, and P. B. Schölkopf, editors, *Advances in Neural Information Processing Systems 16*, pages 659–666. MIT Press, 2004.
- [22] C. Wang, V. Satuluri, and S. Parthasarathy. Local probabilistic models for link prediction. In *Proceedings of the 2007 Seventh IEEE International Conference on Data Mining, ICDM '07*, pages 322–331, Washington, DC, USA, 2007. IEEE Computer Society.
- [23] P. Wang, B. Xu, Y. Wu, and X. Zhou. Link prediction in social networks: the state-of-the-art. *CoRR*, abs/1411.5118, 2014.